

Initialization Heuristics for Greedy Bayesian Network Structure Learning

Walter Perez Urcia Denis Deratani Mauá

Universidade de São Paulo
Instituto de Matemática e Estatística
Departamento de Ciências da Computação

KDMiLE 2015

Contents

- 1 Motivation
- 2 Bayesian Network
- 3 Learning BN
- 4 Initializing Heuristics
- 5 Experiments

Motivation

Car Evaluation Dataset

- Buying price (B): v-high, high, med, low
- Maintain cost (M): v-high, high, med, low
- Doors (D): two, three, four, more
- Persons (P): two, four, more
- Luggage boot (L): small, medium, big
- Safety (S): low, medium, high

Using First-order Logic we have:

- $\forall c, \text{Buying}(c, \text{high}) \rightarrow \text{Doors}(c, \text{four}) \wedge \text{Persons}(c, \text{more})$
- $\forall c, \text{Maintain}(c, \text{low}) \rightarrow \text{Safety}(c, \text{high})$

But, how to represent:


- Half of cars that have four doors have a medium luggage boot
- 15% of cars are low safety, 77% medium safety and 8% high safety

Using probability theory, to represent all possible relations we have:

$$\mathbb{P}(B, M, D, P, L, S)$$

This requires $4 \times 4 \times 4 \times 3 \times 3 \times 3 = 1728$ probabilities hard to estimate, but we can drastically reduce this number by assuming (conditional) independences

For example:



- *D* and *P* are independent given *B*:

$$\mathbb{P}(D, P \mid B) = \mathbb{P}(D \mid B)\mathbb{P}(P \mid B)$$

- *M* and *S* are independent given *D*:

$$\mathbb{P}(M, S \mid D) = \mathbb{P}(M \mid D)\mathbb{P}(S \mid D)$$

⋮

Bayesian Network

A Bayesian Network consists of

- A DAG G over a set of variables X_1, \dots, X_n
- **Markov Property**: Given its parents, every variable is conditionally independent from its non-descendant non-parents
- **Probability constraints**: $\mathbb{P}(X_i = k \mid Pa(X_i) = j) = \theta_{ijk}$

Joint Probability Distribution

There is a unique probability function consistent with a BN:

$$\mathbb{P}(X_1, \dots, X_n) = \prod_{i=1}^n \mathbb{P}(X_i \mid Pa(X_i)) = \prod_{i=1}^n \theta_{ijk}$$

From the previous example:


$$\mathbb{P}(B, M, D, P, L, S) = \mathbb{P}(B)\mathbb{P}(D \mid B)\mathbb{P}(P \mid B)\mathbb{P}(M \mid D)\mathbb{P}(S \mid D, P)\mathbb{P}(L \mid P)$$

This requires

$$4 + (4 \times 4) + (3 \times 4) + (4 \times 4) + (3 \times 4 \times 3) + (3 \times 3) = 93$$


probabilities

Imagine each variable has k values:
 We require k^{33} probabilities without independences.




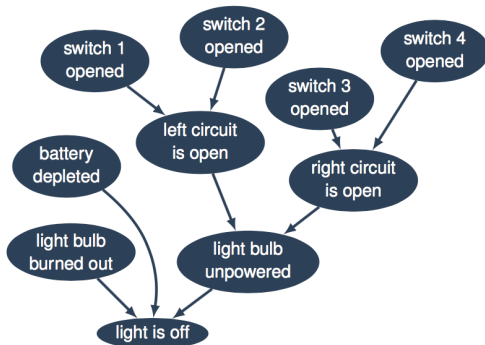
- Elicitation from expert knowledge
- Direct translation
- Learning from data

Elicitation



ANDES: Intelligent Tutoring System to teach Newtonian Physics

Direct Translation



Learning BN

- Age (A): young, adult, old
- Gender (G): male, female
- Education (E): primary, high school, university
- Occupation (O): employee, self-employed
- City size (C): big, small
- Transport (T): private (car), public (bus, train, etc)

| Age | Gender | City | Education | Occupation | Transport | |
|-------|--------|------|-----------|------------|-----------|-------|
| adult | F | | big | high | employee | car |
| adult | M | | small | uni | employee | car |
| adult | F | | big | uni | employee | train |
| young | M | | big | high | self-emp | car |
| adult | M | | big | high | employee | car |
| : | : | | : | : | : | : |

Constraint-based approaches

Perform multiple conditional independence hypothesis testing in order to build a DAG

Score-based approaches

Associate every DAG with a polynomial-time computable score value and search for structure with high score values

Learning as optimization

Given dataset D , select G that maximizes **decomposable** score function:


$$sc(G, D) = F(G) + \psi(N) \times P(G)$$

Learning as optimization


Select G that maximizes **decomposable score function**

$$G^* = \arg \max_{G: G \text{ is a DAG}} sc(G)$$


$$G^* = \arg \max_G \sum_i sc(X_i, Pa(X_i))$$



$$sc(G^*) = -9508.34$$



$$sc(G^*) = -6917.23$$



$$sc(G^*) = -8891.52$$

Greedy Search is a popular approach to find an approximate solution

```


1 GreedySearch( Dataset  $D$  ) : return a BN  $G$ 
2    $G = \text{Initial\_Solution}(X_1, \dots, X_n)$ 
3   For a number of iterations  $K$ 
4      $\text{best\_neighbor} = \text{find\_best\_neighbor}(G)$ 
5     if  $\text{score}(\text{best\_neighbor}) > \text{score}(G)$  then
6        $G = \text{best\_neighbor}$ 
7   Return  $G$ 

```


Different neighborhoods and local moves rise to different methods such as:

- Structure-based
- Equivalence-based
- Order-based


Imagine incumbent solution is




Neighborhood:



Add an edge



Remove an edge



Invert an edge's
direction

Based on the observation that the problem of learning a Bayesian network can be written as

$$G^* = \arg \max_{<} \max_{G \text{ consistent with } <} \sum_{i=1}^n sc(X_i, Pa(X_i))$$

$$G^* = \arg \max_{<} \sum_{i=1}^n \max_{P \subseteq \{X_j < X_i\}} sc(X_i, P)$$

An optimal DAG can be found by maximizing the local scores **independently** given an order of the variables

```

1  OrderBasedGreedySearch( Dataset  $D$  ) : return a BN
2       $L = \text{Get\_Order}(X_1, \dots, X_n)$ 
3      For a number of iterations  $K$ 
4           $\text{current\_sol} = L$ 
5          For each  $i = 1$  to  $n - 1$  do
6               $L_i = \text{swap}(L, i, i + 1)$ 
7              if  $\text{score}(L_i) > \text{score}(\text{current\_sol})$ 
8                   $\text{current\_sol} = L_i$ 
9              if  $\text{score}(\text{current\_sol}) > \text{score}(L)$  then
10                  $L = \text{current\_sol}$ 
11  Return  $\text{network}(L)$ 

```

where $\text{swap}(L, i, i + 1)$ swaps the values $L[i]$ and $L[i + 1]$

Imagine incumbent solution is

$[A, G, E, C, O, T]$

with $sc = -8891.52$

Neighborhood

- $[G, A, E, C, O, T]$, $sc = -7593.82$
- $[A, E, G, C, O, T]$, $sc = -8891.48$
- $[A, G, C, E, O, T]$, $sc = -9149.13$
- $[A, G, E, O, C, T]$, $sc = -6917.23$
- $[A, G, E, C, T, O]$, $sc = -6999.99$

Initializing Heuristics

We propose two different approaches to reduce the space of possible orders:

- DFS-based approach
- FAS-based approach

We can have an upper bound for $sc(G^*)$ by getting $sc(\overline{G})$

$$\overline{G} = \arg \sum_i \max_{Pa(X_i)} sc(X_i, Pa(X_i))$$


Best Parent Set

The parents of a variable X_i in graph \overline{G}

- We can exploit the information provided by \overline{G} and avoid generating orders which are guaranteed sub-optimal
- Assume **best parent set** is unique
- Consider two variables X_i and X_j in \overline{G} , where X_j is parent of X_i , but there is no arc from X_i to X_j
- No optimal ordering can have X_i preceding X_j

The number of these orderings can be much smaller than the full space of orderings

Example

Graph \overline{G}

ABDC

ACDB

ADBC

ADCB

BDAC

BDCA

CDAB

CDBA

DABC

DACB

DBAC

DBCA

DCAB

DCBA

Possible orders from \overline{G}

The algorithm

- Take as input \overline{G} and mark all nodes unvisited
- Start with an empty list L
- While there is an unvisited node
 - Select an unvisited X_i uniformly random
 - Perform a depth-first search (DFS) rooted at X_i and add to L the visited nodes
- Return L

Disadvantage of DFS approach

This approach can be seen as removing edges from \overline{G} such as to make it a DAG and then extract a topological order. But not all edges are **equally relevant** in terms of avoiding poor local maxima.

Estimating the relevance


We can estimate the relevance of an edge $X_j \rightarrow X_i$ by

$$W_{ji} = sc(X_i, Pa^*(X_i)) - sc(X_i, Pa^*(X_i) \setminus \{X_j\})$$

where $Pa^*(X_i)$ represents the **best parent set** for X_i .

We then wish to find a topological ordering of \overline{G} that violates the least cost of edges.

Example



Graph \bar{G}

- C is not very relevant as parent to D
- B is the most relevant parent of D

Min-Cost Feedback Arc Set

Given a weighted directed graph $G = (V, E)$, a set $F \subseteq E$ is called a Min-Cost Feedback Arc Set (min-cost FAS) if every (directed) cycle of G contains at least one edge in F and the sum of weights is minimum.

$$F = \min_{G-F \text{ is a DAG}} \sum_{X_j \rightarrow X_i \in E} W_{ij}$$

Finding FAS F

The following algorithm find an approximate solution:

```

1  MinimumCostFAS( Graph  $G$  ) : Return FAS  $F$ 
2     $F$  = empty set
3    While there is a cycle  $C$  on  $G$  do
4       $W_{min}$  = lowest weight of all edges in  $C$ 
5      For each edge  $(u,v) \in C$  do
6         $W_{uv} = W_{uv} - W_{min}$ 
7        If  $W_{uv} = 0$  add to  $F$ 
8      For each edge in  $F$ , add it to  $G$  if does not build a
        cycle
9    Return  $F$ 

```

The algorithm

- Take the weighted graph \overline{G} with weights W_{ij} as input
- Find min-cost FAS F
- Remove the edges in F from \overline{G}
- Return a topological order from $\overline{G} - F$

Experiments

Configuration

For each dataset considered:

- Limit parent set size to 3
- Perform 1000 runs of Order-based Greedy Search
- At most 100 iterations ($K = 100$)
- Use BIC score
- Find best parent sets by exhaustive search

Datasets

| Dataset | n (#attributes) | N (#instances) | Density of \bar{G} |
|-------------|-----------------|----------------|----------------------|
| Census | 15 | 30168 | 2.85 |
| Letter | 17 | 20000 | 2.41 |
| Image | 20 | 2310 | 2.45 |
| Mushroom | 23 | 8124 | 2.91 |
| Sensors | 25 | 5456 | 3.00 |
| SteelPlates | 28 | 1941 | 2.18 |
| Epigenetics | 30 | 72228 | 1.87 |
| Alarm | 37 | 1000 | 1.98 |
| Spectf | 45 | 267 | 1.76 |
| LungCancer | 57 | 27 | 1.44 |

Table 1: Data sets characteristics

Results

| Dataset | Approach | Best Score | Avg. Initial Score | Avg. Best Score | Avg. It. |
|----------|-----------|-------------------|--|---|-----------------------------------|
| Census | Random | -212186.79 | -213074.18 \pm 558.43 | -212342.26 \pm 174.21 | 7.26 \pm 2.90 |
| | DFS-based | -212190.05 | -212736.80 \pm 379.96 | -212339.83 \pm 152.26 | 5.90 \pm 2.61 |
| | FAS-based | -212191.64 | -212287.99 \pm 92.54 | -212222.12 \pm 70.99 | 3.28 \pm 1.67 |
| Letter | Random | -138652.66 | -139774.54 \pm 413.74 | -139107.13 \pm 329.15 | 6.07 \pm 2.50 |
| | DFS-based | -138652.66 | -139521.38 \pm 396.61 | -138999.84 \pm 310.06 | 5.75 \pm 2.35 |
| | FAS-based | -138652.66 | -139050.43 \pm 70.55 | -139039.26 \pm 87.97 | 2.24 \pm 0.96 |
| Image | Random | -12826.08 | -13017.13 \pm 44.35 | -12924.24 \pm 41.39 | 7.59 \pm 2.71 |
| | DFS-based | -12829.10 | -12999.09 \pm 38.56 | -12921.13 \pm 37.88 | 7.10 \pm 2.47 |
| | FAS-based | -12829.10 | -12930.63 \pm 20.83 | -12882.30 \pm 26.43 | 5.05 \pm 1.72 |
| Mushroom | Random | -55513.38 | -58450.72 \pm 1016.54 | -56563.84 \pm 616.59 | 7.59 \pm 2.76 |
| | DFS-based | -55513.38 | -58367.11 \pm 871.25 | -56472.72 \pm 546.19 | 7.75 \pm 2.58 |
| | FAS-based | -55574.71 | -56450.49 \pm 154.54 | -56198.66 \pm 174.64 | 4.65 \pm 1.63 |
| Sensors | Random | -62062.13 | -63476.33 \pm 265.46 | -62726.60 \pm 251.26 | 9.22 \pm 2.94 |
| | DFS-based | -62083.21 | -63392.60 \pm 255.90 | -62711.50 \pm 257.79 | 9.65 \pm 3.12 |
| | FAS-based | -62074.88 | -62530.26 \pm 133.44 | -62330.94 \pm 121.82 | 5.17 \pm 2.24 |

Results

| Dataset | Approach | Best Score | Avg. Initial Score | Avg. Best Score | Avg. It. |
|-------------|-----------|------------------|--|--|-----------------------------------|
| SteelPlates | Random | -13336.14 | -13566.50 \pm 65.80 | -13429.13 \pm 52.14 | 8.96 \pm 3.43 |
| | DFS-based | -13332.91 | -13572.77 \pm 81.12 | -13432.30 \pm 57.57 | 9.30 \pm 3.38 |
| | FAS-based | -13341.73 | -13485.26 \pm 38.27 | -13397.08 \pm 29.53 | 7.77 \pm 2.24 |
| Epigenetics | Random | -56873.76 | -57722.30 \pm 228.44 | -57357.60 \pm 222.12 | 5.89 \pm 2.67 |
| | DFS-based | -56868.87 | -57615.36 \pm 189.17 | -57308.93 \pm 165.18 | 6.42 \pm 2.47 |
| | FAS-based | -56868.87 | -57660.09 \pm 146.45 | -57379.59 \pm 148.42 | 5.33 \pm 2.28 |
| Alarm | Random | -13218.22 | -13324.52 \pm 30.49 | -13245.43 \pm 15.63 | 10.92 \pm 3.24 |
| | DFS-based | -13217.97 | -13250.72 \pm 17.70 | -13236.71 \pm 12.02 | 4.32 \pm 2.32 |
| | FAS-based | -13220.55 | -13249.77 \pm 2.57 | -13233.98 \pm 6.19 | 6.34 \pm 1.74 |
| Spectf | Random | -8176.81 | -8202.03 \pm 5.23 | -8189.69 \pm 4.65 | 7.20 \pm 2.17 |
| | DFS-based | -8172.37 | -8200.04 \pm 4.08 | -8187.29 \pm 4.91 | 7.86 \pm 2.49 |
| | FAS-based | -8172.51 | -8176.98 \pm 2.01 | -8176.07 \pm 2.05 | 2.27 \pm 1.11 |
| LungCancer | Random | -711.23 | -723.79 \pm 2.69 | -718.03 \pm 2.84 | 5.46 \pm 1.78 |
| | DFS-based | -711.36 | -720.47 \pm 2.51 | -715.29 \pm 1.86 | 5.02 \pm 1.50 |
| | FAS-based | -711.39 | -716.13 \pm 0.89 | -715.67 \pm 1.19 | 2.73 \pm 1.79 |

Results




Figure 1: Census




Figure 2: Letter




Figure 3: Image




Figure 4: Mushroom




Figure 5: Sensors

Results




Figure 6: SteelPlates




Figure 7: Epigenetics




Figure 8: Alarm




Figure 9: Spectf




Figure 10: LungCanc

- The proposed heuristics lead to better solutions on average, and increase the convergence of the search with only a small overhead
- Larger differences for datasets with more variables are expected
- Our proposed techniques could return directed acyclic graphs instead of node orderings to be used for Structure- and Equivalence-based search approaches
- Employ the proposed heuristics in branch-and-bound solvers for finding optimal solutions

Thanks!